

Streaming Multipart Upload to S3 in Java

1. Set Up Your AWS S3 Client

Ensure you have the AWS SDK for Java dependency in your project. For Maven:

```
<dependency>
  <groupId>software.amazon.awssdk</groupId>
  <artifactId>s3</artifactId>
  <version>2.20.18</version> <!-- Replace with the latest version -->
</dependency>
```

Create the S3 client:

```
import software.amazon.awssdk.services.s3.S3Client;
```

```
S3Client s3Client = S3Client.create();
```

```
---
```

2. Initiate a Multipart Upload

Start a multipart upload session and retrieve the uploadId.

```
import software.amazon.awssdk.services.s3.model.CreateMultipartUploadRequest;
import software.amazon.awssdk.services.s3.model.CreateMultipartUploadResponse;
```

```
String bucketName = "your-bucket";
```

```
String keyName = "your-key";
```

```
CreateMultipartUploadRequest createRequest = CreateMultipartUploadRequest.builder()  
    .bucket(bucketName)  
    .key(keyName)  
    .build();
```

```
CreateMultipartUploadResponse createResponse = s3Client.createMultipartUpload(createRequest);  
String uploadId = createResponse.uploadId();
```

```
---
```

3. Upload Parts in Chunks

Stream data in smaller chunks and upload each chunk as a part.

```
import software.amazon.awssdk.services.s3.model.UploadPartRequest;  
import software.amazon.awssdk.services.s3.model.UploadPartResponse;  
import software.amazon.awssdk.core.sync.RequestBody;  
  
import java.io.InputStream;  
import java.util.ArrayList;  
import java.util.List;
```

```
List<CompletedPart> completedParts = new ArrayList<>();  
int partNumber = 1;
```

```
try (InputStream inputStream = getInputStreamFromSource()) { // Replace with your data source
```

```

byte[] buffer = new byte[5 * 1024 * 1024]; // 5 MB buffer

int bytesRead;

while ((bytesRead = inputStream.read(buffer)) != -1) {

    UploadPartRequest uploadPartRequest = UploadPartRequest.builder()

        .bucket(bucketName)

        .key(keyName)

        .uploadId(uploadId)

        .partNumber(partNumber)

        .build();

    UploadPartResponse uploadPartResponse = s3Client.uploadPart(uploadPartRequest,

        RequestBody.fromBytes(buffer, 0, bytesRead));

    completedParts.add(CompletedPart.builder()

        .partNumber(partNumber)

        .eTag(uploadPartResponse.eTag())

        .build());

    partNumber++;

}

}

---
```

4. Complete the Multipart Upload

Once all parts are uploaded, complete the upload by combining the parts.

```
import software.amazon.awssdk.services.s3.model.CompletedMultipartUpload;

import software.amazon.awssdk.services.s3.model.CompletedPart;

import software.amazon.awssdk.services.s3.model.CompleteMultipartUploadRequest;
```

```
CompletedMultipartUpload completedMultipartUpload = CompletedMultipartUpload.builder()

    .parts(completedParts)

    .build();
```

```
CompleteMultipartUploadRequest completeRequest = CompleteMultipartUploadRequest.builder()

    .bucket(bucketName)

    .key(keyName)

    .uploadId(uploadId)

    .multipartUpload(completedMultipartUpload)

    .build();
```

```
s3Client.completeMultipartUpload(completeRequest);
```

```
System.out.println("Multipart upload completed successfully!");
```

```
---
```

5. Abort Multipart Upload on Failure

If any part of the upload fails, abort the multipart upload to avoid leaving incomplete uploads in S3.

```
import software.amazon.awssdk.services.s3.model.AbortMultipartUploadRequest;
```

```
AbortMultipartUploadRequest abortRequest = AbortMultipartUploadRequest.builder()
```

```
.bucket(bucketName)

.key(keyName)

.uploadId(uploadId)

.build();
```

```
s3Client.abortMultipartUpload(abortRequest);

System.out.println("Multipart upload aborted.");
```

6. Helper: Streaming Data Source

Use a streaming data source, such as a database cursor, to feed chunks into the upload process.

```
private InputStream getInputStreamFromSource() {

    // Replace this with your data source logic (e.g., MongoDB cursor or file stream)

    // Example: Use a PipedInputStream to simulate streaming

    return new FileInputStream("path/to/your/large/file");

}
```

Key Points:

- **Chunk Size:** AWS S3 requires each part (except the last) to be at least 5 MB.
- **Concurrency:** To optimize performance, you can upload parts in parallel using `CompletableFuture` or `ExecutorService`.
- **Error Handling:** Ensure robust error handling to retry failed parts or abort the upload if necessary.